


# OCN 750 - HW6

*Maia Kapur*

*September 30, 2015*

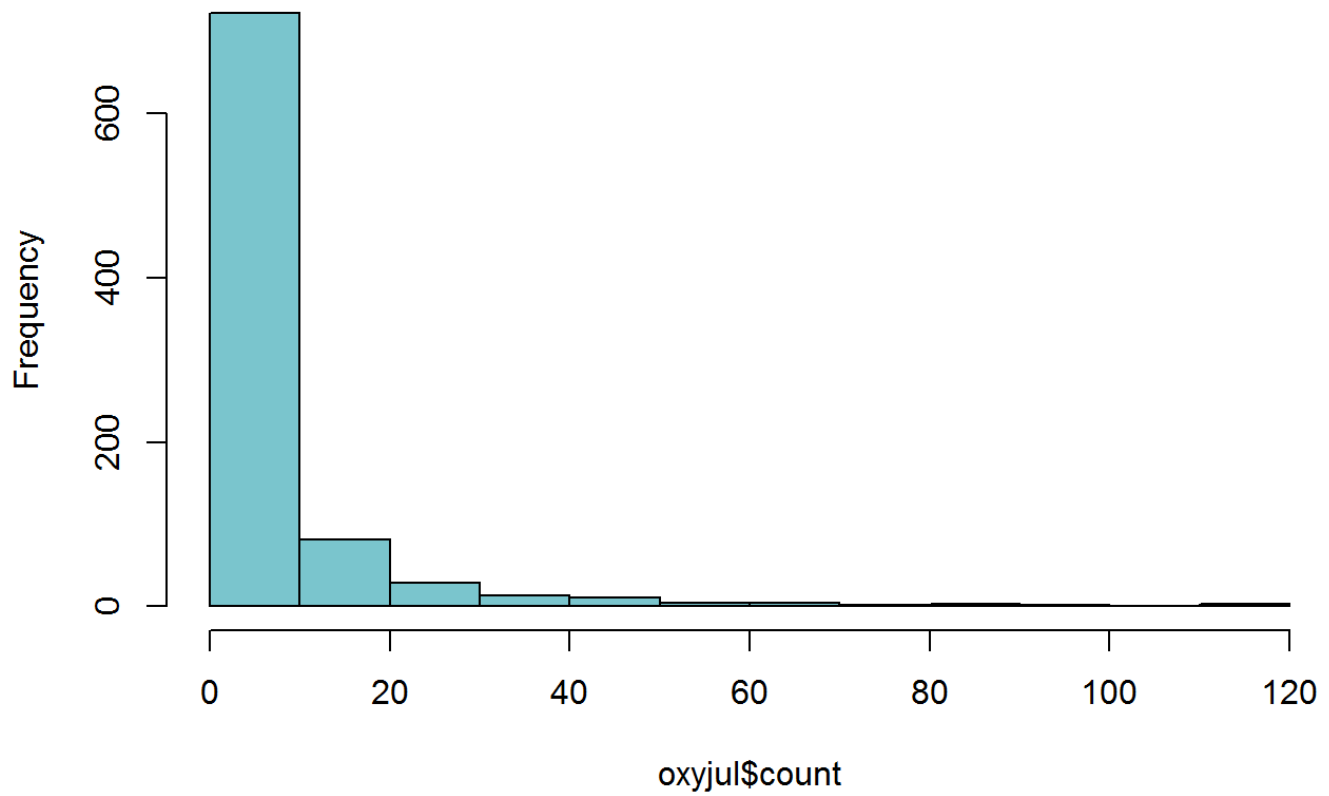
Make some exploratory plots of temperature vs Site, the distribution of all counts, and count vs temperature. What are your thoughts so far?

```
#setwd("~/Dropbox/2015 Fall/OCN 750/hw6")
oxyjul = read.csv("C:/Users/mkapur/Dropbox/2015 Fall/OCN 750/hw6/oxyjulis_subset.csv")
plot(Temperature ~ Site.x, data = oxyjul, main = "temperature vs. site", pch = 19, col =
"darkorange1", col.main = "darkorange1")
```




```
hist(oxyjul$count, col = "cadetblue3", main = "distribution of counts", col.main = "cadet
blue3")
```

## distribution of counts



```
plot(count ~ Temperature, data = oxyjul, pch = 19, col = "lavenderblush3", main = "Oxyjul  
is count vs temp", col.main = "lavenderblush3")
```

## Oxyjulis count vs temp



*##The counts are highly skewed with many zeros, possibly inflated. Even without the zeros there doesn't seem to be a big trend between count and temperature. There seem to be some slight trends amongst the two groups of sites. So...this is tricky.*

This is count data, so starting with a Poisson GLM seems like a good idea. There are many samples per site, so we should include Site as a predictor (make sure to make it a factor!). Make a Poisson GLM with temperature and Site as predictors, and quantify how overdispersed the data is.

```
as.factor(oxyjul$Site.x) ##coerce the site vals to factor
```

```
site.modp = glm(count ~ Site.x + Temperature, data = oxyjul, family = "poisson") ##generate your poisson glm  
summary(site.modp)
```

```
##
## Call:
## glm(formula = count ~ Site.x + Temperature, family = "poisson",
##      data = oxyjul)
##
## Deviance Residuals:
##      Min        1Q    Median        3Q        Max
## -5.0642  -3.3243  -2.4741  -0.0862   21.1793
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  0.255220    0.112622   2.266   0.0234 *
## Site.x        0.022628    0.001255  18.028  <2e-16 ***
## Temperature  0.077774    0.006784  11.464  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for poisson family taken to be 1)
##
##      Null deviance: 13937  on 872  degrees of freedom
## Residual deviance: 13438  on 870  degrees of freedom
## AIC: 15211
##
## Number of Fisher Scoring iterations: 6
```

```
overdis = function(model) {
  sum(residuals(model, type = "pearson")^2)/(length(model$y) - length(model$coefficients))
} ##create the overdispersion function
overdis(site.modp) ##this parameter is very high, meaning overdispersion is low
```


```
## [1] 29.30449
```

Based on the shape of the data, it seems like a good idea to consider zero inflation. But we'll want to compare a zero-inflated model to a non-inflated model, so first fit a negative binomial GLM with temperature and site as predictors. Plot the fitted effects, and report likelihood ratio tests for the predictors.


```
##fit the neg-binomial GLM
require(MASS)
nb.mod = glm.nb(count ~ Site.x + Temperature, data = oxyjul)

##Plot of fitted effects
require(effects)
plot(allEffects(nb.mod))
```

**Site.x effect plot**



**Temperature effect plot**



```
##Anova LRT for predictors
```

```
require(car)
```

```
Anova(nb.mod)
```

```
## Analysis of Deviance Table (Type II tests)
```

```
##
```

```
## Response: count
```

```
##          LR Chisq Df Pr(>Chisq)
```

```
## Site.x      9.8753  1  0.001675 **
```

```
## Temperature  2.6996  1  0.100371
```

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

What's the estimate for theta for the negative binomial?

```
theta = overdis(nb.mod)
```

```
theta ##less than 1.5
```

```
## [1] 1.193846
```

What does this model say about the relationship between abundance and


temperature?

```
##The relationship looks nearly linear though the CIs are wide. So Temperature is somewhat positively correlated with the model, which isn't greatly overdispersed.
```



Plot the deviance residuals vs 1) the fitted values, 2) vs the predictors. Does a linear model for temperature (on the link scale) seem reasonable?

```
##plot of deviance residuals vs fitted values  
plot(resid(nb.mod, family = "pearson") ~ fitted(nb.mod), main = "deviance residuals vs fitted values", pch = 19, col = "indianred2", col.main = "indianred2")
```

deviance residuals vs fitted values



```
##plot of deviance residuals vs. predictors for two params  
par(mfrow = c(2,1))  
plot(resid(nb.mod, family = "pearson") ~ oxyjul$Site.x, main = "residuals vs site", pch = 19, col = "indianred3", col.main = "indianred3", ylab = "pearson residuals", xlab = "site")  
  
plot(resid(nb.mod, family = "pearson") ~ oxyjul$Temperature, main = "residuals vs temperature", pch = 19, col = "indianred4", col.main = "indianred4", ylab = "pearson residuals", xlab = "temp")
```



```
dev.off()
```

```
## null device
##          1
```

*##A linear model doesn't make sense to me. There seems to be a big asymptote/gap in the data and the residuals show a slight negative trend as predictors increase.*

For now let's keep it somewhat simple and make 5 models to compare: 1) the negative binomial model with temperature + site;

```
##See above.
##AIC for nb.mod. Store this for later comparison
nb.mod.aic = AIC(nb.mod)
```

2) a zero-inflated poisson with temperature + site as predictors only for the count model;

```
require(pscl)
as.data.frame(oxyjul)
modz = zeroinfl(count ~ Temperature + Site.x | 1, data = oxyjul, dist = "poisson")
modz.aic = AIC(modz)
```

3) a zero-inflated negative binomial with temperature + site as predictors only for the count model;

```
modz.nb = zeroinfl(count ~ Temperature + Site.x | 1, data = oxyjul, dist = "negbin")
modz.nb.aic = AIC(modz.nb)
```

4) a zero-inflated poisson with temperature + site as predictors both for the count model and also for the zero-inflation (binomial) model;

```
##similar to number 2, but now you have the predictors on both sides of the line, indicating them both for the poisson and binomial parts
modz.both = zeroinfl(count ~ Temperature + Site.x | Temperature + Site.x, data = oxyjul, dist = "poisson")
modz.both.aic = AIC(modz.both)
```

5) a zero-inflated negative binomial with temperature + site as predictors both for the count model and also for the zero-inflation (binomial) model.

```
##add in the "as factor" so R knows to separate the sites
modz.nb.both = zeroinfl(count ~ Temperature + as.factor(Site.x) | Temperature + as.factor(Site.x), data = oxyjul, dist = "negbin")
modz.nb.both.aic = AIC(modz.nb.both)
```

Calculate AIC for all five models. Which is the 'best' model (lowest AIC)? What does it mean that this is the best model, compared to the other models?

```
aics = cbind(modz.aic, modz.nb.aic, modz.both.aic, modz.nb.both.aic)
names(aics[1:4]) = c("MODZ.Poisson", "MODZ.NegBin", "MODZ.Poisson.Both", "MODZ.NegBin.Both")
min(aics) ##occurs at negative binomial model with predictors in both the count model and binomial model.
```

```
## [1] 4139.507
```


```
##this estimates the model's quality, as an estimate of the information lost by using a given model. It also accounts for the complexity of the model in a sort of "tradeoff" between good-fit and overfitting. This model does the best job in minimizing unexplained data while being fairly simple. Normally we look for AICs under 10, I think.
```

What is the effect(s) of temperature in the best model? How do you interpret this result, statistically and biologically, compared to the negative binomial model



(i.e. non-inflated)?

```
summary(nb.mod) ##the negative binomial model only suggests significance for "site" upon  
the abundance  
summary(modz.nb.both) ##whereas the best model shows significance in both the count model  
and zero-inflated binomial model. The effect is less significant and quite small in the f  
ormer, however.  
plot(count ~ Temperature, data = oxyjul)
```



```
##This suggests that the extra zeros may have masked the effect of temperature, which has  
as strong and significant effect without them.
```

Do you notice anything a little funny about the standard errors for the model coefficients for the best model?

```
##They seem really small
```

For model #5 in the above list, let's visualize the fitted effect of temperature on the probability of getting an 'extra' zero. Because the model has a factor for Site, that means the different sites will have different intercepts (but the same slope for temperature). So pick four sites and use `curve()` to plot the fitted logistic curve for those sites, all on the same plot, with the x-axis having the same range as the

range of temperature in the dataset.

```
##plot presence/absence of counts vs temp. This is what was done for possum absence vs.
stags in lec 10.
plot(absence ~ Temperature, data = oxyjul, pch = 20, col = "gray 73", main = "Oxyjul. abs
ence vs temperature", sub = "fitted zero-inflation lines for four sites", col.main = "med
iumpurple1", col.sub = "mediumpurple")

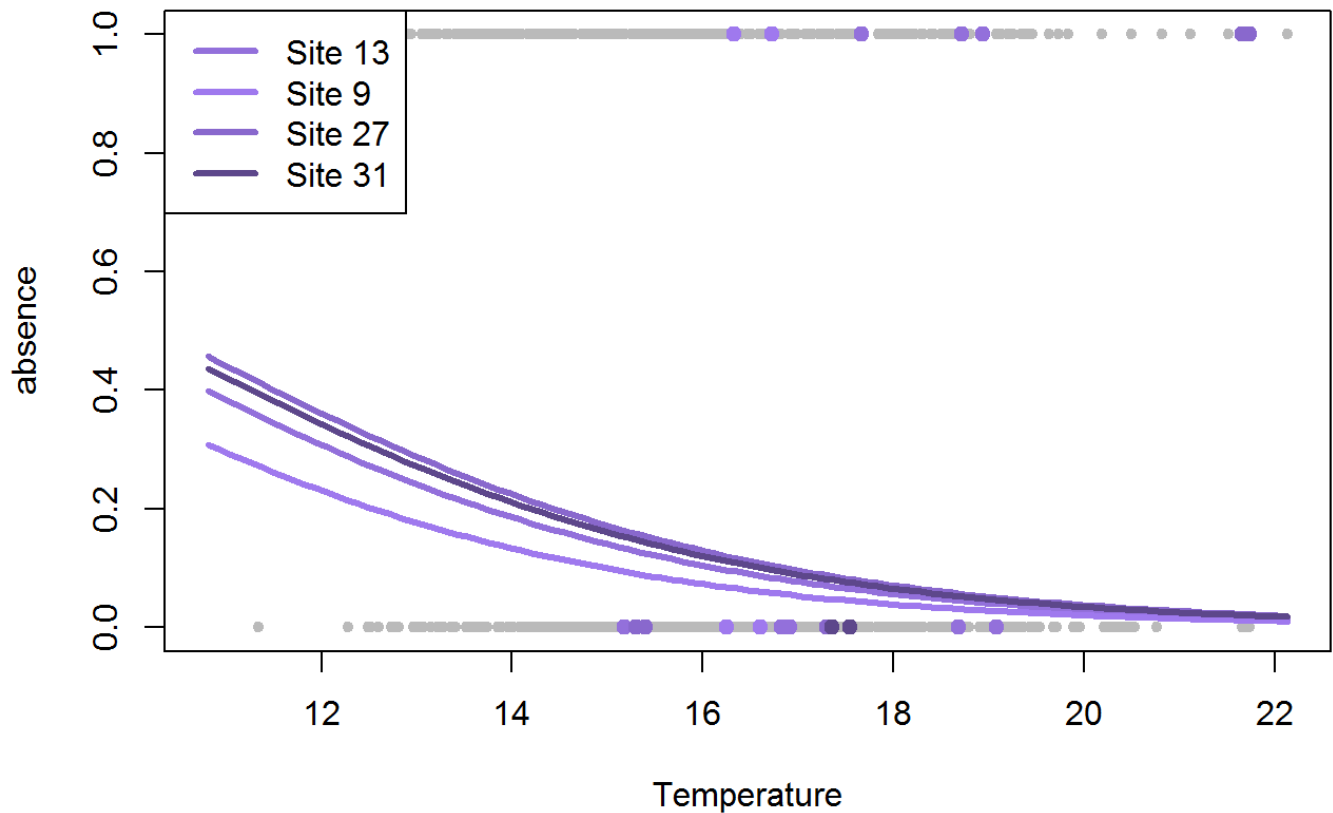
##highlight the sites' points, for reference
with(subset(oxyjul, Site.x == 13), points(count ~ Temperature, pch = 19, col = 'me
diumpurple'))
with(subset(oxyjul, Site.x == 9), points(count ~ Temperature, pch = 19, col = "me
diumpurple2"))
with(subset(oxyjul, Site.x == 27), points(count ~ Temperature, pch = 19, col = "me
diumpurple3"))
with(subset(oxyjul, Site.x == 31), points(count ~ Temperature, pch = 19, col = "me
diumpurple4"))

##function for logistic curve
logistic = function(x) {exp(x)/(1+exp(x))}

##the offsets are pretty small for each, and the slope for temperature is minimal
curve(logistic(coef(modz.nb.both)["count_(Intercept)"] + coef(modz.nb.both)["count_as.fac
tor(Site.x)13"] + coef(modz.nb.both)["zero_Temperature"] * x), add = TRUE, col = "mediump
urple",lwd = 3)
curve(logistic(coef(modz.nb.both)["count_(Intercept)"] + coef(modz.nb.both)["count_as.fac
tor(Site.x)9"] + coef(modz.nb.both)["zero_Temperature"] * x), add = TRUE, col = "mediumpu
rple2",lwd = 3)
curve(logistic(coef(modz.nb.both)["count_(Intercept)"] + coef(modz.nb.both)["count_as.fac
tor(Site.x)27"] + coef(modz.nb.both)["zero_Temperature"] * x), add = TRUE, col = "mediump
urple3",lwd = 3)
curve(logistic(coef(modz.nb.both)["count_(Intercept)"] + coef(modz.nb.both)["count_as.fac
tor(Site.x)31"] + coef(modz.nb.both)["zero_Temperature"] * x), add = TRUE, col = "mediump
urple4", lwd = 3)

##add a Legend
legend('topleft', lty = 1, lwd = 3, col = c('mediumpurple', 'mediumpurple2', 'med
iumpurple3', 'mediumpurple4'), legend = c('Site 13', 'Site 9', 'Site 27', 'Site 31'))
```

## Oxyjul. absence vs temperature



fitted zero-inflation lines for four sites

Roughly how much does temperature change the proportion of extra zeros?